

In the context of areal units the general inferential issues are the following:

- Of course, how much smoothing is appropriate is not readily defined. In particular, for model-based smoothers such as we describe below, it is not evident what the extent of smoothing is, or how to control it. Specification of a utility function for smoothing (as attempted in Stern and Cressie, 1999) would help to address these questions but does not seem to be considered in practice..

- 73

Equations (4.5) and (4.6) could simultaneously hold only in trivial cases, so the two mean specifications are *incompatible*. Thus we can say that $f(y_1|y_2)$ and $f(y_2|y_1)$ are incompatible with regard to determining $p(y_1, y_2)$. We do not examine conditions for compatibility of conditional distributions here, although there has been considerable work in this area (see, e.g., Arnold and Strauss, 1991, and references therein).

Another point is that $p(y_1 \dots, y_n)$ may be improper even if $p(y_i|y_j, j \neq i)$ is proper for all i . As an elementary illustration, consider $p(y_1, y_2) \propto \exp[-\frac{1}{2}(y_1 - y_2)^2]$. Evidently $p(y_1|y_2)$ is $N(y_2, 1)$ and $p(y_2|y_1)$ is $N(y_1, 1)$, but $p(y_1, y_2)$ is improper. Casella and George (1992) provide a similar example in a bivariate exponential (instead of normal) setting.

Brook's Lemma notes that

$$\begin{aligned} p(y_1, \dots, y_n) &= \frac{p(y_1|y_2, \dots, y_n)}{p(y_{10}|y_2, \dots, y_n)} \cdot \frac{p(y_2|y_{10}, y_3, \dots, y_n)}{p(y_{20}|y_{10}, y_3, \dots, y_n)} \\ &\quad \dots \frac{p(y_n|y_{10}, \dots, y_{n-1,0})}{p(y_{n0}|y_{10}, \dots, y_{n-1,0})} \cdot p(y_{10}, \dots, y_{n0}), \end{aligned} \quad (4.7)$$

an identity which is easily checked (Exercise 1). Here, $\mathbf{y}_0 = (y_{10}, \dots, y_{n0})'$ is any fixed point in the support of $p(y_1, \dots, y_n)$. Hence $p(y_1, \dots, y_n)$ is determined by the full conditional distributions, since apart from the constant $p(y_{10}, \dots, y_{n0})$ they are the only objects appearing on the right-hand side of (4.7). Hence the joint distribution is determined up to a proportionality constant. If $p(y_1, \dots, y_n)$ is improper then this is, of course, the best we can do; if $p(y_1, \dots, y_n)$ is proper then the fact that it integrates to 1 determines the constant. Perhaps most important is the constructive nature of (4.7): we can create $p(y_1, \dots, y_n)$ simply by calculating the product of ratios. For more on this point, see Exercise 2.

When the number of areal units is very large (say, a regular grid of pixels associated with an image or a large number of small geographic regions), we do not seek to write down the joint distribution of the Y_i . Rather we prefer to work (and model) exclusively with the n corresponding full conditional distributions. In fact, from a spatial perspective, we would imagine that the full conditional distribution for Y_i would be more "local," that is, it should really depend only upon the neighbors of cell i . Adopting some definition of a neighbor structure (e.g., the one setting $W_{ij} = 1$ or 0 depending on whether i and j are adjacent or not), let ∂_i denote the set of neighbors of cell i .

Next suppose we specify a set of full conditional distributions for the Y_i such that

$$p(y_i|y_j, j \neq i) = p(y_i|y_j, j \in \partial_i) \quad (4.8)$$

A critical question to ask is whether a specification such as (4.8) uniquely determines a joint distribution for Y_1, \dots, Y_n . That is, we do not need to see the explicit form of this distribution. We merely want to be assured that if, for example, we implement a Gibbs sampler (see Subsection 5.3.1) to simulate realizations from the joint distribution, that there is indeed a unique stationary distribution for this sampler.

The notion of using *local* specification to determine a joint (or global) distribution in the form (4.8) is referred to as a *Markov random field* (MRF). There is by now a substantial literature in this area, with Besag (1974) being a good place to start. Geman and Geman (1984) provide the next critical step in the evolution, while Kaiser and Cressie (2000) offer a current view and provide further references. See also Rue and Held (2005) and references therein.

A critical definition in this regard is that of a *clique*. A clique is a set of cells (equivalently, indices) such that each element is a neighbor of every other element. With n cells, depending upon the definition of the neighbor structure, cliques can possibly be of size 1, 2, and so on up to size n . A *potential function* (or simply *potential*) of order k is a function of k arguments that is exchangeable in these arguments. The arguments of the potential would be the values taken by variables associated with the cells for a clique of size k . For continuous

Referring to (4.12), we may write the entire system of random variables as

$$(I - B)\mathbf{Y} = \boldsymbol{\epsilon}. \quad (4.18)$$

When $p(\mathbf{y})$ is proper we can appeal to standard multivariate normal distribution theory to interpret the entries in $\Sigma_{\mathbf{y}}^{-1}$. For example, $1/(\Sigma_{\mathbf{y}}^{-1})_{ii} = \text{Var}(Y_i|Y_j, j \neq i)$. Of course with $\Sigma_{\mathbf{y}}^{-1} = D^{-1}(I - B)$, $(\Sigma_{\mathbf{y}}^{-1})_{ii} = 1/\tau_i^2$ providing immediate agreement with (4.12). But also, if $(\Sigma_{\mathbf{y}}^{-1})_{ij} = 0$, then Y_i and Y_j are conditionally independent given $Y_k, k \neq i, j$, a fact you are asked to show in Exercise 10. Hence if any $b_{ij} = 0$, we have conditional independence for that pair of variables. Connecting b_{ij} to w_{ij} shows that the choice of neighbor structure implies an associated collection of conditional independences. With first-order neighbor structure, all we are asserting is a spatial illustration of the local Markov property (Whittaker, 1990, p. 68).

We also note that in suitable contexts it may be appropriate to think of \mathbf{Y}_i as a vector of dependent areal unit measurements or, in the context of random effects, as a vector of dependent random effects associated with an areal unit. This leads to the specification of multivariate conditionally autoregressive (MCAR) models, which is the subject of Section 10.1. From a somewhat different perspective, \mathbf{Y}_i might arise as $(Y_{i1}, \dots, Y_{iT})^T$ where Y_{it} is the measurement associated with areal unit i at time t , $t = 1, \dots, T$. Now we would of course think in terms of spatiotemporal modeling for Y_{it} . This is the subject of Section 11.7.

Finally, Gaussian Markov random fields can introduce proximities more general than those that we have discussed here. In particular, working with a regular lattice, there is much scope for further theoretical development. For instance, Rue and Held (2005, p. 114) describe the derivation of the following model weights based on the forward difference analogue of penalizing the derivatives of a surface used to specify the thin plate spline. They consider 12 neighbors of a given point. The north, east, south, and west neighbors

$$p(y_1, y_2, \dots, y_n; \psi) \propto \exp(\psi \sum_{i,j} w_{ij} 1(y_i = y_j)) \quad (4.21)$$

4.4 Simultaneous autoregressive (SAR) models

$$\mathbf{Y} \sim N\left(\mathbf{0}, (I - B)^{-1} \tilde{D} ((I - B)^{-1})^T\right). \quad (4.22)$$

Alternatively, W can be replaced by \widetilde{W} where now, for each i , the i th row has been normalized to sum to 1. That is, $(\widetilde{W})_{ij} = w_{ij}/w_{i+}$. Again, \widetilde{W} is not symmetric, but it is row stochastic, i.e., $\widetilde{W}\mathbf{1} = \mathbf{1}$. If we set $B = \alpha\widetilde{W}$, α is called a *spatial autocorrelation parameter* and, were W a contiguity matrix, now $Y_i = \alpha \sum_j Y_j I(j \in \partial_i)/w_{i+} + \epsilon_i$. With a very regular grid the w_{i+} will all be essentially the same and thus α will be a multiple of ρ . But, perhaps more importantly, with \widetilde{W} row stochastic the eigenvalues of \widetilde{W} are all less than or equal to 1 (i.e., $\max |\lambda_i| = 1$). Thus $I - \alpha\widetilde{W}$ will be nonsingular if $\alpha \in (-1, 1)$, justifying referring to α as an autocorrelation parameter; see Exercise 9.

4.4.1 CAR versus SAR models

Cressie (1993, pp. 408–10) credits Brook (1964) with being the first to make a distinction between the CAR and SAR models, and offers a comparison of the two. To begin with, we may note from (4.13) and (4.22) that, under propriety, the two forms are equivalent if and only if

$$(I - B)^{-1}D = (I - \tilde{B})^{-1}\tilde{D}((I - \tilde{B})^{-1})^T,$$

where we use the tilde to indicate matrices in the SAR model. Cressie then shows that any SAR model can be represented as a CAR model (since D is diagonal, we can straightforwardly solve for B), but gives a counterexample to prove that the converse is not true. Since all SAR models are proper while we routinely employ improper CAR models, it is not surprising that the latter is a larger class.

For the “proper” CAR and SAR models that include spatial correlation parameters ρ , Wall (2004) shows that the correlations between neighboring regions implied by these two models can be rather different; in particular, the first-order neighbor correlations increase at a slower rate as a function of ρ in the CAR model than they do for the SAR model. (As an aside, she notes that these correlations are not even monotone for $\rho < 0$, another reason to avoid negative spatial correlation parameters.) Also, correlations among pairs can switch in nonintuitive ways. For example, when working with the adjacency relationships generated by the lower 48 contiguous U.S. states, she finds that when $\rho = .49$ in the CAR model, $Corr(Alabama, Florida) = .20$ and $Corr(Alabama, Georgia) = .16$. But when ρ increases to .975, we instead get $Corr(Alabama, Florida) = .65$ and $Corr(Alabama, Georgia) = .67$, a slight reversal in ordering.

4.4.2 STAR models

In the literature SAR models have frequently been extended to handle spatiotemporal data. The idea is that in working with proximity matrices, we can define neighbors in time as well as in space. Figure 4.4 shows a simple illustration with 9 areal units, 3 temporal units for each areal unit yielding $i = 1, \dots, 9$, $t = 1, 2, 3$, labeled as indicated.

The measurements Y_{it} are spatially associated at each fixed t . But also, we might seek to associate, say, Y_{i2} with Y_{i1} and Y_{i3} . Suppose we write Y as the 27×1 vector with the first nine entries at $t = 1$, the second nine at $t = 2$, and the last nine at $t = 3$. Also let $W_S = \text{BlockDiag}(W_1, W_1, W_1)$, where

$$W_1 = \begin{pmatrix} 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \end{pmatrix}.$$

Then W_S provides a spatial contiguity matrix for the Y 's. Similarly, let $W_T = \begin{pmatrix} 0 & W_2 & 0 \\ W_2 & 0 & W_2 \\ 0 & W_2 & 0 \end{pmatrix}$, where $W_2 = I_{9 \times 9}$. Then W_T provides a *temporal* contiguity matrix for the Y 's. But then, in our SAR model we can define $B = \rho_s W_S + \rho_t W_T$. In fact, we can

Consider, for example, the SAT exam scores data from the lower 48 contiguous states of the U.S. and the District of Columbia. We can construct this map using the `maps` and `maptools` packages. The idea is to create a sequence of data structures that will eventually produce an adjacency matrix. We execute the following commands

The object `"usa.state"` is returned by the `"map()"` function. Next, we extract the state ID's from the `"state"` database and use them to construct a `"SpatialPolygon"` object called `"usa.poly"`, which is then converted to a neighborhood object `"usa.nb"` using the `spdep` function `poly2nb`. Finally the `nb2mat` function (also in `spdep`) produces the adjacency matrix. The option `style="B"` produces the basic binary coding. Therefore, `usa.adj.mat` produced above is a 49×49 matrix whose (i, j) -th entry is equal to 1 if i is a neighbor of j and 0 otherwise. All diagonal entries are 0. The option `style="W"` produces a row-normalized adjacency matrix. Other options are available and left to the reader to explore.

```
> mn.county = map("county", "minnesota", fill=TRUE, plot=FALSE)
> county.ID <- sapply(strsplit(mn.county$names, ","), function(x) x[2])
> mn.poly = map2SpatialPolygons(mn.county, IDs=county.ID)
> mn.nb = poly2nb(mn.poly)
> mn.adj.mat = nb2mat(mn.nb, style="B")
```

```
> mn.region.id <- attr(mn.nb, "region.id")
> winona.neighbors.index = mn.nb[[match("winona", mn.region.id)]]
> winona.neighbors = rownames(mn.adj.mat[winona.neighbors.index,])
> winona.neighbors
[1] "fillmore" "houston"  "olmsted"  "wabasha"
```

One could also create adjacency matrices from external shapefiles by executing

06/30/2019 - RS00000000000000000000001096850 (Rachel Ulrich) - Hierarchical Modeling and Analysis for Spatial Data, Second Edition

The first step produces a `SpatialPolygonsDataFrame` object `nc.sids`, while the second step produces the region IDs and stores them in `rn`. The third step uses these region IDs to produce an `nb` object by directly reading from the GAL file. We next use a Freeman-Tukey transformation to produce the transformed rates and append them to the `nc.sids` object.

We wish to regress these rates on the non-white birth rates over the same period. This variable is available as `NWBIR79` in the `nc.sids` object. We will use the Freeman-Tukey transformed birth rates:

Maximum likelihood estimation of (4.23), which has the likelihood in (4.24), can be carried out using the `errorsarlm()` or, equivalently, the `spautolm()` function in `spdep`. These functions produce the same output. Below we demonstrate the latter. We first create a `listw` object using the 0-1 adjacency structure

Note that the `zero.policy=TRUE` is required here because the county shapefile in `spdep` lists two counties, Dare and Hyde, as having zero neighbors. These counties are situated in coastal North Carolina and are adjacent to substantial bodies of water. These two counties can be identified as

We now estimate the SAR model in (4.23) using

Call:

Residuals:

Regions with no neighbours included:

Coefficients:

06/30/2019 - RS0000000000000000000000001096850 (Rachel Ulrich) - Hierarchical Modeling and Analysis for Spatial Data, Second Edition

the regression coefficients do not change substantially and are still significant, the spatial autocorrelation parameter is no longer statistically significant (p -value is approximately 0.25257). The result from the CAR model (obtained by setting `family="CAR"`) is also very similar.

4.6 Exercises

1. Verify Brook's Lemma, equation (4.7).
- 2.(a) To appreciate how Brook's Lemma works, suppose Y_1 and Y_2 are both binary variables, and that their joint distribution is defined through conditional logit models. That is,

$$\log \frac{P(Y_1 = 1|Y_2)}{P(Y_1 = 0|Y_2)} = \alpha_0 + \alpha_1 Y_2 \quad \text{and} \quad \log \frac{P(Y_2 = 1|Y_1)}{P(Y_2 = 0|Y_1)} = \beta_0 + \beta_1 Y_1 .$$

Obtain the joint distribution of Y_1 and Y_2 .

- (b) This result can be straightforwardly extended to the case of more than two variables, but the details become increasingly clumsy. Illustrate this issue in the case of *three* binary variables, Y_1 , Y_2 , and Y_3 .
3. Returning to (4.13) and (4.14), let $B = ((b_{ij}))$ be an $n \times n$ matrix with positive elements; that is, $b_{ij} > 0$, $\sum_j b_{ij} \leq 1$ for all i , and $\sum_j b_{ij} < 1$ for at least one i . Let $D = \text{Diag}(\tau_i^2)$ be a diagonal matrix with positive elements τ_i^2 such that $D^{-1}(I - B)$ is symmetric; that is, $b_{ij}/\tau_i^2 = b_{ji}/\tau_j^2$, for all i, j . Show that $D^{-1}(I - B)$ is positive definite.
4. Looking again at (4.13), obtain a simple sufficient condition on B such that the CAR specification with precision matrix $D^{-1}(I - B)$ is a pairwise difference specification, as in (4.16).
5. Show that, for W symmetric, $\Sigma_{\mathbf{y}}^{-1} = D_w - \rho W$ is positive definite (thus resolving the impropriety in (4.15)) if $\rho \in (1/\lambda_{(1)}, 1/\lambda_{(n)})$, where $\lambda_{(1)} < \lambda_{(2)} < \dots < \lambda_{(n)}$ are the ordered eigenvalues of $D_w^{-1/2} W D_w^{-1/2}$.
6. Show that if all entries in W are nonnegative and $D_w - \rho W$ is positive definite with $0 < \rho < 1$, then all entries in $(D_w - \rho W)^{-1}$ are nonnegative.
7. Under a proper CAR model for \mathbf{Y} , i.e., with $\Sigma_{\mathbf{y}} = D_w - \rho W$, obtain the correlation and covariance between Y_i and Y_j .
8. Obtain the joint distribution, up to normalizing constant, for (Y_1, Y_2, \dots, Y_n) under (4.20). Hint: You might try to guess it but Brook's Lemma can be used as well.
9. Recalling the SAR formulation using the scaled adjacency matrix \widetilde{W} just below (4.22), prove that $I - \alpha \widetilde{W}$ will be nonsingular if $\alpha \in (-1, 1)$, so that α may be sensibly referred to as an "autocorrelation parameter."
10. In the setting of Subsection 4.3.1, if $(\Sigma_{\mathbf{y}}^{-1})_{ij} = 0$, then show that Y_i and Y_j are conditionally independent given $Y_k, k \neq i, j$.
11. The file `www.biostat.umn.edu/~brad/data/state-sat.dat` gives the 1999 state average SAT data (part of which is mapped in Figure 4.1).
 - (a) Use the `spautolm` function to fit the SAR model of Section 4.4, taking the verbal SAT score as the response Y and the percent of eligible students taking the exam in each state as the covariate X . Do this analysis twice: first using binary weights and then using row-normalized weights. Is the analysis sensitive to these choices of weights? Is knowing X helpful in explaining Y ?
 - (b) Using the `maps` library in R, draw choropleth maps similar to Figure 4.1 of both the fitted verbal SAT scores and the spatial residuals from the SAR model. Is there

